

```

import torch
import torch.nn as nn
import torch.nn.functional as F

block_size = 8
batch_size = 4
data_text = "hello world"
vocab = sorted(list(set(data_text)))
vocab_size = len(vocab)
device = 'cuda' if torch.cuda.is_available() else 'cpu'
epochs = 1000
embed_size = 32

stoi = {ch: i for i, ch in enumerate(vocab)}
itos = {i: ch for ch, i in stoi.items()}
encode = lambda s: [stoi[c] for c in s]
decode = lambda l: ''.join([itos[i] for i in l])

data = torch.tensor(encode(data_text), dtype=torch.long).to(device)

def get_batch():
    ix = torch.randint(len(data) - block_size, (batch_size,))
    x = torch.stack([data[i:i + block_size] for i in ix])
    y = torch.stack([data[i + 1:i + block_size + 1] for i in ix])
    return x, y

class TinyLLM(nn.Module):
    def __init__(self):
        super().__init__()
        self.embed = nn.Embedding(vocab_size, embed_size)
        self.fc = nn.Linear(embed_size * block_size, vocab_size)

    def forward(self, x):
        x = self.embed(x)                # (B, T, C)
        x = x.view(x.size(0), -1)        # flatten: (B, T*C)
        return self.fc(x)                # (B, vocab_size)

# --- Model Setup ---
model = TinyLLM().to(device)
optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)

# --- Training Loop ---
for epoch in range(epochs):
    x, y = get_batch()
    logits = model(x)
    loss = F.cross_entropy(logits, y[:, -1]) # Only last token predicted
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    if epoch % 100 == 0:
        print(f"Epoch {epoch}, Loss: {loss.item():.4f}")

➡ Epoch 0, Loss: 1.9726
Epoch 100, Loss: 0.0096
Epoch 200, Loss: 0.0043
Epoch 300, Loss: 0.0025
Epoch 400, Loss: 0.0018
Epoch 500, Loss: 0.0013
Epoch 600, Loss: 0.0010
Epoch 700, Loss: 0.0008
Epoch 800, Loss: 0.0006
Epoch 900, Loss: 0.0005

# --- Generation Function ---
def generate(model, start='h', max_new_tokens=20):
    model.eval()
    idx = torch.tensor([encode(start)], dtype=torch.long).to(device)

    for _ in range(max_new_tokens):
        idx_cond = idx[:, -block_size:]

```

```
# Padding if needed
if idx_cond.shape[1] < block_size:
    padding = torch.zeros(1, block_size - idx_cond.shape[1], dtype=torch.long).to(device)
    idx_cond = torch.cat([padding, idx_cond], dim=1)

logits = model(idx_cond)
probs = F.softmax(logits, dim=-1)
next_id = torch.multinomial(probs, num_samples=1)
idx = torch.cat((idx, next_id), dim=1)

return decode(idx[0].tolist())

# --- Generate text after training ---
print("Generated text:", generate(model))

🔄 Generated text: hhed ddrldwlh wdrld
```

Start coding or [generate](#) with AI.